

Hierarchical Clustering

Prof. Marco Alvarez, Computer Science
University of Rhode Island

Hierarchical clustering

- ▶ Unsupervised learning technique that groups observations into nested clusters
- ▶ Key features:
 - does not require a predefined number of clusters
 - creates a hierarchy of clusters, often visualized as a dendrogram
 - can be agglomerative (bottom-up), or divisive (top-down)
 - widely used in various domains (gene expression, phylogenetic trees, community detection, etc.)

Agglomerative vs divisive

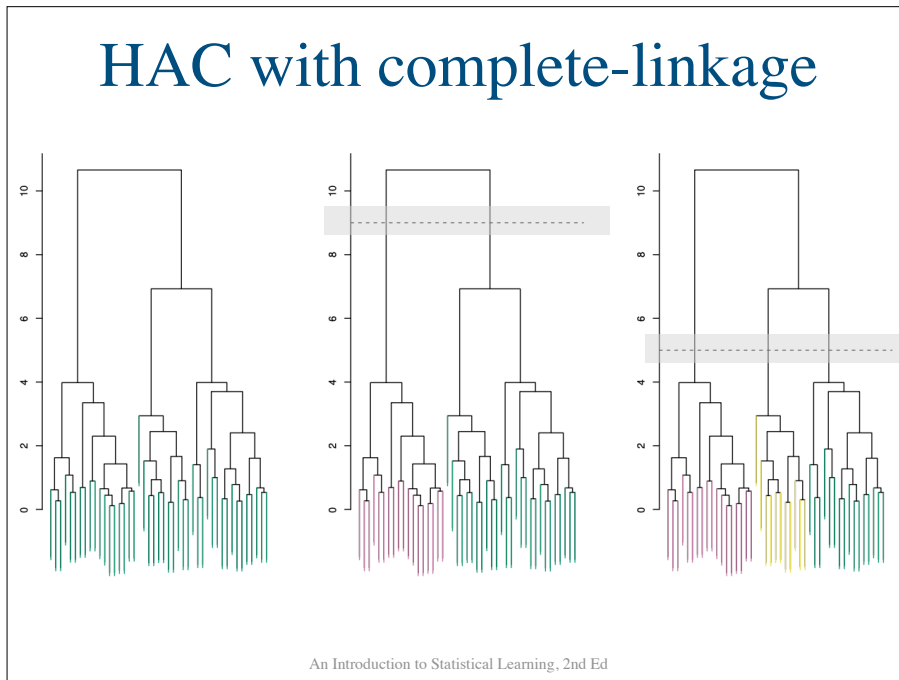

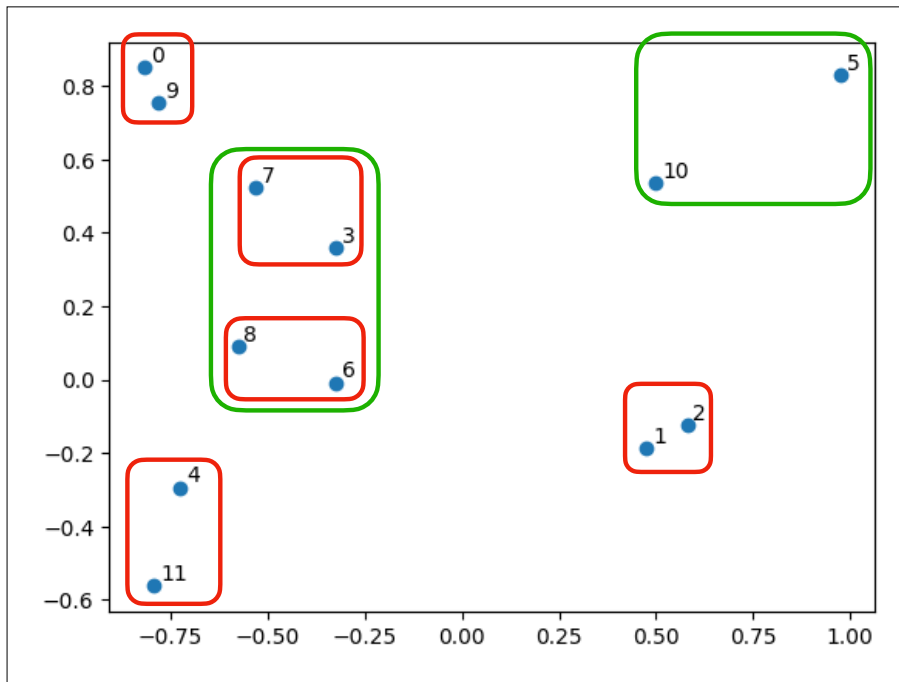
- ▶ Agglomerative approach (**bottom-up**)
 - start with observations as singleton clusters and progressively merge them until a final single cluster is created
 - most common form
- ▶ Divisive approach (**top-down**)
 - start with all observations on a single cluster and progressively split the clusters until all observations are singleton clusters

HAC algorithm

- ▶ Initially each observation is considered a cluster
- ▶ Repeat
 - **merge**: find the closest pair of clusters and merge it into a single cluster
 - if all observations are in a single cluster stop

Distance Metrics:
- Euclidean distance
- Manhattan distance
- Minkowski distance
- Hamming distance
- Cosine similarity
- Correlation coefficient

Linkage Criteria:
- Single-linkage
- Complete-linkage
- Average-linkage
- Centroid-linkage
- Ward's method



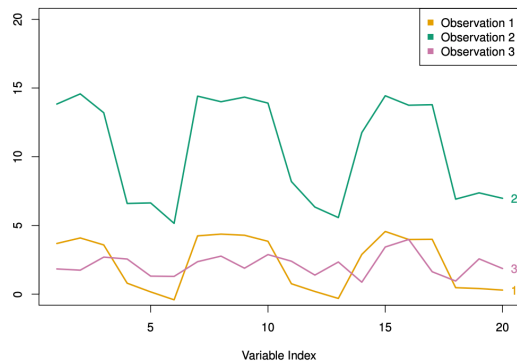
Pairwise metrics

See the [Pairwise metrics, Affinities and Kernels](#) section of the user guide for further details.

| | |
|--|--|
| <code>metrics.pairwise.additive_chi2_kernel(X[, Y])</code> | Compute the additive chi-squared kernel between observations in X and Y. |
| <code>metrics.pairwise.chi2_kernel(X[, Y, gamma])</code> | Compute the exponential chi-squared kernel between X and Y. |
| <code>metrics.pairwise.cosine_similarity(X[, Y, ...])</code> | Compute cosine similarity between samples in X and Y. |
| <code>metrics.pairwise.cosine_distances(X[, Y])</code> | Compute cosine distance between samples in X and Y. |
| <code>metrics.pairwise.distance_metrics()</code> | Valid metrics for pairwise_distances. |
| <code>metrics.pairwise.euclidean_distances(X[, Y, ...])</code> | Compute the euclidean distance matrix between each pair from a vector array X and Y. |
| <code>metrics.pairwise.haversine_distances(X[, Y])</code> | Compute the Haversine distance between samples in X and Y. |
| <code>metrics.pairwise.kernel_metrics()</code> | Valid metrics for pairwise_kernels. |
| <code>metrics.pairwise.laplacian_kernel(X[, Y, gamma])</code> | Compute the laplacian kernel between X and Y. |
| <code>metrics.pairwise.linear_kernel(X[, Y, ...])</code> | Compute the linear kernel between X and Y. |
| <code>metrics.pairwise.manhattan_distances(X[, Y, ...])</code> | Compute the L1 distances between the vectors in X and Y. |
| <code>metrics.pairwise.nan_euclidean_distances(X)</code> | Calculate the euclidean distances in the presence of missing values. |
| <code>metrics.pairwise.pairwise_kernels(X[, Y, ...])</code> | Compute the kernel between arrays X and optional array Y. |
| <code>metrics.pairwise.polynomial_kernel(X[, Y, ...])</code> | Compute the polynomial kernel between X and Y. |
| <code>metrics.pairwise.rbf_kernel(X[, Y, gamma])</code> | Compute the rbf (gaussian) kernel between X and Y. |
| <code>metrics.pairwise.sigmoid_kernel(X[, Y, ...])</code> | Compute the sigmoid kernel between X and Y. |
| <code>metrics.pairwise.paired_euclidean_distances(X, Y)</code> | Compute the paired euclidean distances between X and Y. |
| <code>metrics.pairwise.paired_manhattan_distances(X, Y)</code> | Compute the paired L1 distances between X and Y. |
| <code>metrics.pairwise.paired_cosine_distances(X, Y)</code> | Compute the paired cosine distances between X and Y. |
| <code>metrics.pairwise.paired_distances(X, Y, *[, ...])</code> | Compute the paired distances between X and Y. |
| <code>metrics.pairwise_distances(X[, Y, metric, ...])</code> | Compute the distance matrix from a vector array X and optional Y. |
| <code>metrics.pairwise_distances_argmin(X, Y, *[, ...])</code> | Compute minimum distances between one point and a set of points. |
| <code>metrics.pairwise_distances_argmin_min(X, Y, *)</code> | Compute minimum distances between one point and a set of points. |
| <code>metrics.pairwise_distances_chunked(X[, Y, ...])</code> | Generate a distance matrix chunk by chunk with optional reduction. |

What kind of distance to consider?

- We have been using Euclidean distances, however, the choice is **very important**

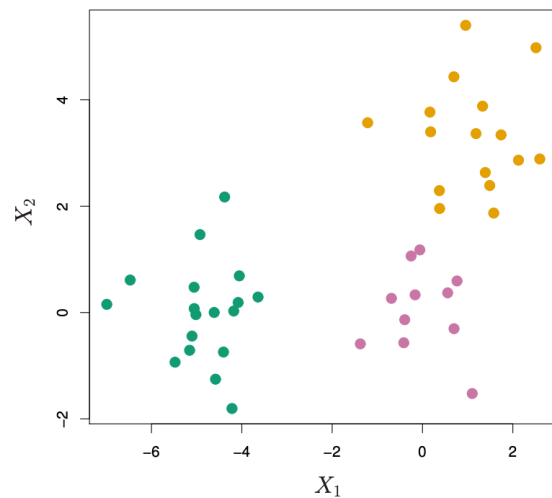


An Introduction to Statistical Learning, 2nd Ed

Distances between clusters

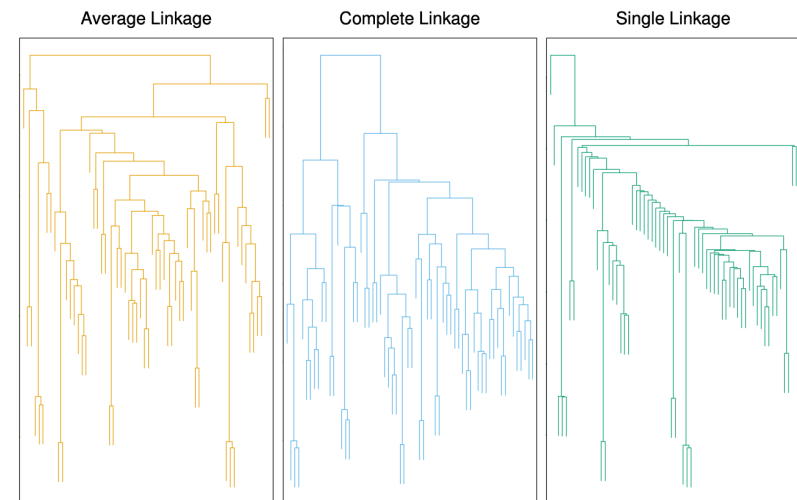
| | |
|-------------------|--|
| Single-linkage: | $D(A, B) = \arg \min_{\mathbf{x}_1 \in A, \mathbf{x}_2 \in B} d(\mathbf{x}_1, \mathbf{x}_2)$ |
| Complete-linkage: | $D(A, B) = \arg \max_{\mathbf{x}_1 \in A, \mathbf{x}_2 \in B} d(\mathbf{x}_1, \mathbf{x}_2)$ |
| Average-linkage: | $D(A, B) = \frac{1}{ A B } \sum_{\mathbf{x}_1 \in A, \mathbf{x}_2 \in B} d(\mathbf{x}_1, \mathbf{x}_2)$ |
| Centroid-linkage: | $D(A, B) = \ \mu_A - \mu_B\ _2^2$ |
| Ward's method: | Minimizes variance within clusters |

Consider this data



An Introduction to Statistical Learning, 2nd Ed

Example



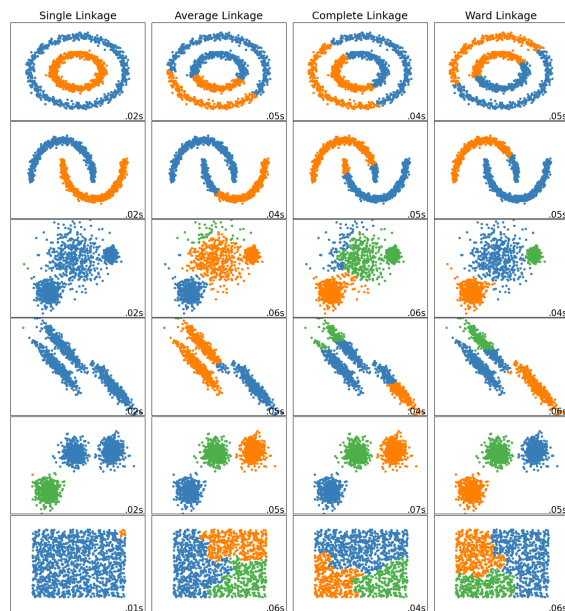
An Introduction to Statistical Learning, 2nd Ed

Limitations and considerations

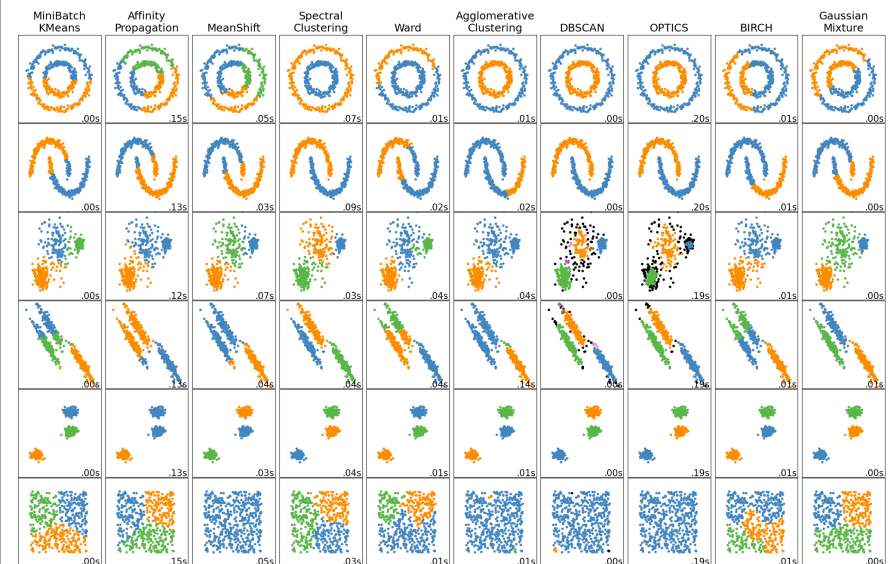
- ▶ **Computational complexity**
 - can be computationally expensive for large datasets
- ▶ **Sensitivity to noise and outliers**
 - can be sensitive to noise and outliers in the data
- ▶ **Choice of distance metric and linkage criterion**
 - distance metric and linkage criterion can significantly affect the results
- ▶ **Optimal number of clusters**
 - no definitive method for determining the optimal number of clusters

HAC Notebook

<https://colab.research.google.com/drive/13ENliESw1VT7IWNXeB0Q-7tyMuHPppuM>



From scikit-learn



From scikit-learn