

Dimensionality Reduction

Prof. Marco Alvarez, Computer Science
University of Rhode Island

High dimensional data

- Prevalent in many areas of machine learning and data science
- Examples:
 - image data, for instance a single 1000 by 1000 RGB image has 3 million dimensions
 - text data, in natural language processing text is often represented in high-dimensional spaces, transformers typically embed tokens into 768 or more dimensions
 - genomic data, gene expression data often has thousands of dimensions.
 - time series from sensor data or financial data
 - audio signals, especially when converted to spectrograms
 - network traffic data
 - ...

Dimensionality reduction

- Fundamentally ...
 - unsupervised learning algorithms for extracting latent structure (potentially **low dimensional**) from **high-dimensional** data
 - can range from simple feature selection to complex nonlinear transformations
- Examples
 - PCA, Kernel PCA, t-SNE, autoencoders, matrix factorization
- Given $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ with $\mathbf{x}_i \in \mathbb{R}^d$, find a representation $\mathcal{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_n\}$ with $\mathbf{z}_i \in \mathbb{R}^{d'}$, with $d' < d$
 - certain properties should be preserved (e.g., variance, distances, neighborhood structure)


Dimensionality reduction

- Why?
 - visualization (2D or 3D)
 - preprocessing data before machine learning
 - focusing on important features/patterns
 - more efficient training
 - removing noise and redundant information
 - data compression

$$\begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_n^T \end{bmatrix}_{n \times d} \Rightarrow \begin{bmatrix} \mathbf{z}_1^T \\ \vdots \\ \mathbf{z}_n^T \end{bmatrix}_{n \times d'}$$

Project 2D data onto a line


$$\begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_n^T \end{bmatrix}_{11 \times 2}$$




$$\begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_n^T \end{bmatrix}_{11 \times 1}$$



Which line is better? which is worst?




Maximize variance



Idea: minimize sum of square distances to the line (or maximize the sum of squares of the scalar projections)

Variance



“biased” estimator default in numpy

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

Different data, same variance




$$\text{x-variance} = \frac{2^2 + 0^2 + 2^2}{3} = 8/3$$

$$\text{y-variance} = \frac{1^2 + 0^2 + 1^2}{3} = 2/3$$

Figure credit: <https://serrano.academy/unsupervised-learning/>


Covariance



$$\text{cov}(\mathbf{x}, \mathbf{y}) = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

Figure credit: <https://serrano.academy/unsupervised-learning/>

Covariance matrix



Every element Σ_{ij} of the covariance matrix is the covariance between column i and column j from the data matrix

Eigenvectors and eigenvalues

- ▶ The decomposition of a **square matrix A** into eigenvalues and eigenvectors is known as **eigen decomposition**
- for **real symmetric matrices** eigenvectors can be chosen real and orthonormal


$$\mathbf{A} = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^T \quad \mathbf{A} \mathbf{v} = \lambda \mathbf{v}$$

columns of \mathbf{V} are the eigenvectors of \mathbf{A} and $\mathbf{\Lambda}$ is a diagonal matrix whose entries are the eigenvalues of \mathbf{A}

Principal Component Analysis (PCA)

Goal of PCA

- Find projections of the data onto directions that maximize variance
 - these directions are **orthogonal** to each other



<https://www.dataschool.io/15-hours-of-expert-machine-learning-videos/>

PCA approach

- Input: data matrix $X_{d \times n}$
 - center** the data (subtract the mean)
 - calculate the **covariance** matrix $\frac{1}{n}XX^T$
 - compute **eigendecomposition** $V\Lambda V^T$ of the covariance matrix
 - sort** the eigenvectors by eigenvalues in decreasing order
- Output
 - sorted orthonormal eigenvectors V and eigenvalues Λ

eigenvectors can then be used for projecting the data into lower dimensions (XV)

Remarks

- The **larger** the eigenvalue, the **more important** the corresponding eigenvector
 - that's why we **sort** eigenvalues (and corresponding eigenvectors) in **decreasing order**
- All eigenvalues of a positive semidefinite matrix are **non-negative**
 - covariance matrix** is always symmetric and p.s.d.
- For dimensionality reduction, we can ignore eigenvectors associated with smaller eigenvalues

Explained variance

- ▶ Each eigenvalue corresponds to the amount of variance explained by its associated eigenvector
 - **explained variance** is often presented as a **percentage**, i.e., eigenvalues divided by the total sum of eigenvalues
- ▶ The sum of percentages of the top-k principal components is usually referred to as the **“cumulative explained variance”**
 - often used to select how many components to keep for a reduced dataset

Explained variance

PC	Eigenvalue	Variance (%)	Cumulative Variance
1st	23.31800	59.072%	59.072%
2nd	7.01200	17.764%	76.835%
3rd	4.61800	11.699%	88.534%
4th	1.98100	5.018%	93.553%
5th	1.00100	2.536%	96.089%
6th	0.82100	2.080%	98.168%
7th	0.64100	1.624%	99.792%
8th	0.03100	0.079%	99.871%
9th	0.02900	0.073%	99.944%
10th	0.02200	0.056%	100.000%

PCA Notebooks

<https://colab.research.google.com/drive/1MzPdVsJi8gUxhwiXFcKA8RJsw8DY1OhE>

<https://colab.research.google.com/drive/1r7JPdmmWS11y12WVOMi0GMhlDM9s37GI>