

# CSC 461: Machine Learning

Fall 2024

## Supervised learning

Prof. Marco Alvarez, Computer Science  
University of Rhode Island

## Definition

- Paradigm in ML where a model learns from labeled data
  - mapping inputs to corresponding outputs by minimizing a predefined “loss function”
- Key components
  - data instance:  $(x, y), x \in \mathcal{X}, y \in \mathcal{Y}$ 
    - input space  $\mathcal{X}$
    - output space  $\mathcal{Y}$
  - training data:  $\{(x_1, y_1), \dots, (x_n, y_n)\} \subseteq \mathcal{X} \times \mathcal{Y}$
  - model (hypothesis):  $h: \mathcal{X} \mapsto \mathcal{Y}, h \in \mathcal{H}$ 
    - hypothesis space  $\mathcal{H}$

## Dataset

- Observations are independently drawn from a **joint distribution** of inputs and outputs

$$\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$$


- Input space  $\mathcal{X}$

- numerical — continuous (age, income, ...)
- categorical (gender, product type, ...)
- text (customer reviews, documents, ...)
- images (photos, medical images, ...)
- audio (speech, music, ...)


$$(x_i, y_i) \sim P$$

**unknown**

## Example: MNIST dataset



## Example: MNIST dataset



Image

```
[[ 0.  1.  5. 11. 15.  4.  0.  0.]
 [ 0.  8. 16. 13.  6.  2.  0.  0.]
 [ 0. 11.  7.  0.  0.  0.  0.  0.]
 [ 0. 11. 16. 16. 11.  2.  0.  0.]
 [ 0.  0.  4.  4.  5. 12.  3.  0.]
 [ 0.  0.  0.  0.  0.  5. 11.  0.]
 [ 0.  0.  1.  6.  0. 10. 11.  0.]
 [ 0.  0.  2. 12. 16. 15.  2.  0.]]
```

Matrix representation

```
[ 0.  1.  5. 11. 15.  4.  0.  0.  0.  8. 16. ... 11.  0.  0.  0.  2. 12. 16. 15.  2.  0.]
```

Vector representation

$\mathcal{X} =$

$\mathcal{Y} =$

## Major tasks

## Types of supervised learning


### ► Regression

- continuous output:  $\mathcal{Y} \subseteq \mathbb{R}$
- examples: predicting house prices, forecasting stock prices

### ► Classification


- discrete output:  $\mathcal{Y} = \{1, \dots, K\}$ 
  - binary classification:  $K = 2$
  - multi-class classification:  $K > 2$
- examples, spam vs not spam (binary), disease present (binary), handwritten digits recognition (multi-class)

## Regression



$\mathcal{X} =$

$\mathcal{Y} =$




$\mathcal{X} =$

$\mathcal{Y} =$


## Binary classification

```
array([[ 0.24277092,  0.89098144], array([[0],  
[-0.57961074,  0.50618765], [1],  
[ 0.24259841,  0.12209649], [1],  
[ 1.68348295, -0.10059047], [1],  
[ 2.00696736, -0.79306007], [1],  
[ 1.56891881,  0.30515286], [0],  
[ 0.1314049 , -0.35704446], [1],  
[ 2.14017386,  0.33933491], [1],  
[-1.03087047,  1.52609949], [0],  
[-0.38504321,  1.24209655], [0],  
[-1.20252537,  0.56167652], [0],  
[ 0.08590311,  0.68265315], [1],  
[ 0.88074085, -0.11759523], [1],  
[ 0.32558238,  0.4181143 ], [1],  
[-0.74202798,  0.68847344]])
```

## Binary classification

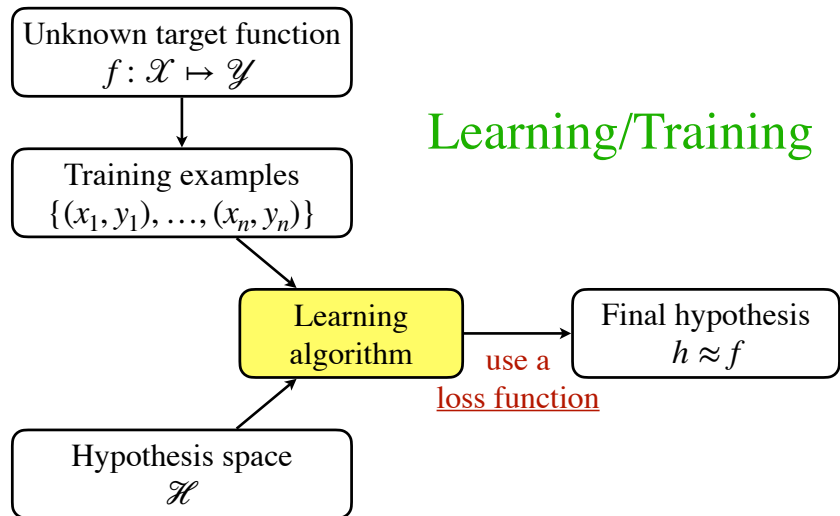


## Multi-class classification

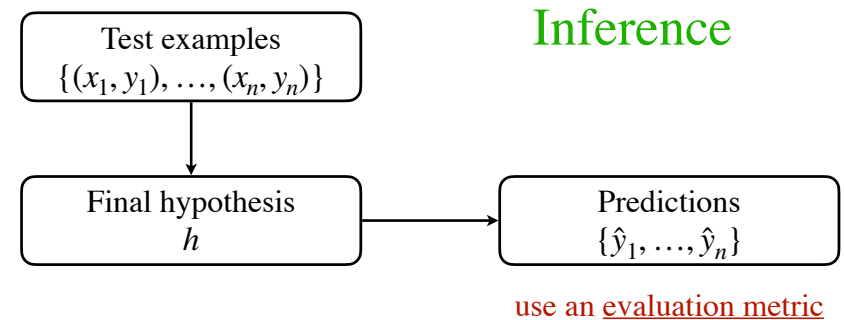


Learning

## Supervised learning setup



## Supervised learning setup



## Loss functions and evaluation

### Loss functions

#### ► Purpose

- guide the learning process during training/learning

#### ► Characteristics

- differentiable (in most cases — neural networks)
- optimized during training/learning
- reflect model performance on individual examples or batches of examples

## 0/1 loss

$$l_{0/1}(h, x_i, y_i) = I(h(x_i) \neq y_i)$$

indicator function

Prediction	Target
5	5
1	9
2	2
7	7
8	0
0	0
0	8
3	3
6	6
4	4

## Squared loss (L2 loss)

$$l_{sq}(h, x_i, y_i) = (h(x_i) - y_i)^2$$

Prediction	Target
1.2	1.4
2.3	2.3
1.1	1.2
3.4	4.1
2.3	2.5
1.1	1.1
2.5	2.6
3.1	3.2
1.7	1.8
2.3	2.3

positive loss,  
penalizes  
big mistakes

## Absolute loss

$$l_{abs}(h, x_i, y_i) = |h(x_i) - y_i|$$

Prediction	Target
1.2	1.4
2.3	2.3
1.1	1.2
3.4	4.1
2.3	2.5
1.1	1.1
2.5	2.6
3.1	3.2
1.7	1.8
2.3	2.3

## Evaluation metrics

### ▸ Purpose

- assess model performance after training

### ▸ Characteristics

- may not be differentiable
- used to compare different models/hypotheses or report final performance
- often more interpretable
- reflect performance on an entire dataset

### ▸ Examples

- accuracy, f1-score, precision, recall, mean absolute error, R-squared

# More formally ...

## From statistical learning theory

### Expected Loss

- theoretical average loss over **all possible data points**
- including those not in our training set

$$\mathbb{E} \left[ l(h, x_i, y_i) \right]_{(x_i, y_i) \sim P}$$

We cannot calculate this term, but we can approximate it

### Empirical Loss

- the average loss calculated on the **training data**
- what we can actually measure during training/learning

$$\frac{1}{n} \sum_{i=1}^n l(h, x_i, y_i)$$

## Learning formulation

### Given:

- training set:  $\{(x_1, y_1), \dots, (x_n, y_n)\}, x_i \in \mathcal{X}, y_i \in \mathcal{Y}$ ,

### Objective:

- find a function  $h : \mathcal{X} \rightarrow \mathcal{Y}, h \in \mathcal{H}$  that minimizes the **empirical loss**  $L$

$$L = \frac{1}{n} \sum_{i=1}^n l(h, x_i, y_i)$$

As the size of our training dataset increases, the **empirical loss** tends to approach the **expected loss**

## Challenges

### Overfitting

- model performs well on training data but poorly on unseen data
- solutions include: regularization, data augmentation

### Underfitting

- model is too simple to capture underlying patterns
- solutions include: increase model complexity

### Imbalanced Datasets

- one class may be overrepresented.
- solutions include: resampling, class weighting

### Curse of Dimensionality

- too many features can cause models to perform poorly
- solutions include: dimensionality reduction, using more data