

CSC 461: Machine Learning

Fall 2024


Linear regression

Prof. Marco Alvarez, Computer Science
University of Rhode Island

Preliminaries

Continuous output



- Certain applications require the prediction of continuous values



<https://towardsdatascience.com/stock-prediction-in-python-b66555171a2>

Linear functions

- Assumes the output y is a linear function of the input x
- can use the function to make predictions, very simple approach, e.g. linear regression




$$y = wx + b \quad w, x, b \in \mathbb{R}$$

slope intercept

Linear functions

► What if we have d features?

- $\mathbf{w}, \mathbf{x} \in \mathbb{R}^d$
- $b \in \mathbb{R}$



$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

hypothesis weights bias

The weights and bias are the model **parameters** which define the hypothesis and are used to make predictions

Linear regression

Alternative notation

► Augmented vectors

- incorporate bias into the weight vector
 - $\mathbf{w} = [b, w_1, w_2, \dots, w_d]^T$
- augment input vector with 1
 - $\mathbf{x} = [1, x_1, x_2, \dots, x_d]^T$
- simplified equation
 - $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$

\mathbf{x}_1	\mathbf{x}_2	\mathbf{y}
0.5	0.1	0.25
0.3	0.9	0.5
0.3	0.875	1.15
0.45	0.15	2.13
...



\mathbf{x}_0	\mathbf{x}_1	\mathbf{x}_2	\mathbf{y}
1	0.5	0.1	0.25
1	0.3	0.9	0.5
1	0.3	0.875	1.15
1	0.45	0.15	2.13
...

Hypothesis space $\mathcal{H} = \{h_{\mathbf{w}} : h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}, \mathbf{w} \in \mathbb{R}^{d+1}\}$

Linear regression

► Fundamental supervised learning algorithm

- used for predicting continuous target variables
- assumes a linear relationship between input features and the target

► Applications

- financial forecasting (e.g., stock prices, economic indicators)
- environmental science (e.g., climate change predictions)
- marketing (e.g., sales forecasting)
- real estate (e.g., house price prediction)

Linear regression

▸ Data

- $\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$
- $\mathbf{x}^{(i)} \in \mathbb{R}^{d+1}, y^{(i)} \in \mathbb{R}$

▸ Model

- $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$


▸ Squared Loss (L2)

- $l_{sq}(h, \mathbf{x}, y) = (h(\mathbf{x}) - y)^2$
- $L(h, \mathcal{D}) = \frac{1}{n} \sum_{i=1}^n l_{sq}(h, \mathbf{x}^{(i)}, y^{(i)})$

Residuals and MSE loss

▸ Residual


- difference between predicted and actual value



Want to find a **linear function** with **small residuals**

Closed-form solution

min vs. argmin



$$f(x) = 6x^2 + 4x + 5$$

$\min f(x)?$

$\arg \min_{x} f(x)?$

Given a function, it represents the input value(s) that produce the function's minimum output value

Normal equations

- Analytical solution to minimize the loss function

$$\mathbf{w}^* = \arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (h(\mathbf{x}^{(i)}) - y^{(i)})^2$$



$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$

Using matrix notation

$$\mathbf{y} = \mathbf{X}\mathbf{w} \quad \begin{bmatrix} (\mathbf{x}^{(1)})^T \\ \vdots \\ (\mathbf{x}^{(n)})^T \end{bmatrix} \begin{bmatrix} w_0 \\ \vdots \\ w_d \end{bmatrix} \approx \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

Using matrix notation

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$




$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

Minimizing the loss function

- Set the gradient to zero

- solve for \mathbf{w}

continuous,
differentiable,
convex



$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

Minimizing the loss function

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla E_{\text{in}}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = \mathbf{0}$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y} \quad \text{where} \quad \mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$

\mathbf{X}^\dagger is the 'pseudo-inverse' of \mathbf{X}

There are other methods for finding the optimal solution
e.g. gradient descent, MLE

<http://work.caltech.edu/slides/slides03.pdf>

Show me the code

```
# generate random data
Xtr = np.hstack((np.ones((100, 1)), np.random.rand(100, 4)))
Ytr = np.random.rand(100, 1)
Xte = np.hstack((np.ones((10, 1)), np.random.rand(10, 4)))
Yte = np.random.rand(10, 1)
```

```
w = np.linalg.pinv(Xtr) @ Ytr
pred = Xte @ w
```

or

```
np.linalg.inv(Xtr.T @ Xtr) @ Xtr.T @ Ytr
pred = Xte @ w
```

or

```
reg = LinearRegression().fit(Xtr, Ytr)
pred = reg.predict(Xte)
```

```
loss = np.mean((pred-Yte) ** 2)
print(loss)
```

Conclusion

▸ Linear regression: simple yet powerful

- foundation for many advanced ML techniques
- serves as a baseline for more complex models

▸ Computational Complexity

- becomes inefficient for large datasets or high-dimensional data
- time complexity: $O(nd^2 + d^3)$

▸ Alternatives

- mini-batch gradient descent
- regularization techniques (Ridge, Lasso)

▸ Limitations and Considerations

- assumes linearity between features and target
- sensitive to outliers