

CSC 461: Machine Learning

Fall 2024

Feature transformation

Prof. Marco Alvarez, Computer Science
University of Rhode Island

Data transformation

Definition

- converting or mapping data from its raw form into a format that better suits the learning algorithms
- crucial step in the machine learning pipeline that can significantly improve model performance

Why?

- meet algorithm assumptions
- handle non-linear relationships
- scale features appropriately
- handle different data types

Scaling

Min-max scaling

$$\tilde{x} = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

Standardization

$$\tilde{x} = \frac{x - \mu}{\sigma}$$

```
[[ -91  25 -85 -36]
 [  13  23  56  57]
 [  13 -92 -27 -100]
 [ -60  64  15 -84]
 [   0  39 -46 -12]]
```

`MinMaxScaler()`

```
[[0.      0.75      0.      0.40764331]
 [1.      0.73717949 1.      1.      ]
 [1.      0.      0.41134752 0.      ]
 [0.29807692 1.      0.70921986 0.10191083]
 [0.875     0.83974359 0.27659574 0.56050955]]
```

`StandardScaler()`

```
[[-1.54757414  0.24479055 -1.38466458 -0.01789141]
 [ 0.89102753  0.20770107  1.50346717  1.64600941]
 [ 0.89102753 -1.92494389 -0.19663876 -1.16294143]
 [-0.82068325  0.96803537  0.66365581 -0.87667893]
 [ 0.58620232  0.5044169  -0.58581963  0.41150235]]
```

Categorical transformations

One-hot encoding

- converts categorical variables into binary vectors
- each category becomes a column, only one column contains 1, the rest are 0

```
['red']      [[0. 0. 1.]
['green']    [0. 1. 0.]
['blue']     [1. 0. 0.]
['red']      [0. 0. 1.]
['green']    [0. 1. 0.]
['blue']     [1. 0. 0.]
['red']      [0. 0. 1.]
['green']    [0. 1. 0.]
['blue']     [1. 0. 0.]
```

Label encoding

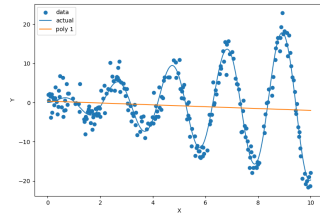
- converts categorical values into numerical by assigning a unique integer to each category

```
['red']      [2]
['green']    [1]
['blue']     [0]
['red']      [2]
['green']    [1]
['blue']     [0]
['red']      [2]
['green']    [1]
['blue']     [0]
```

Non-linear transformations

► Motivation

- data is not always “linear”



$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix} \xrightarrow{\Phi} \mathbf{z} = \begin{bmatrix} \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_{\tilde{d}}(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} z_1 \\ \vdots \\ z_{\tilde{d}} \end{bmatrix}$$

input space $\mathcal{X} = \mathbb{R}^d$

feature space $\mathcal{Z} = \mathbb{R}^{\tilde{d}}$

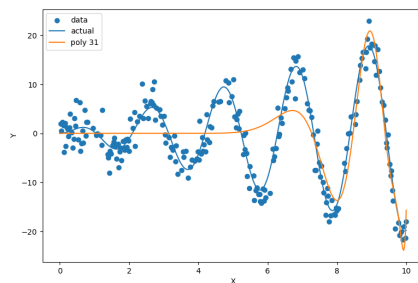
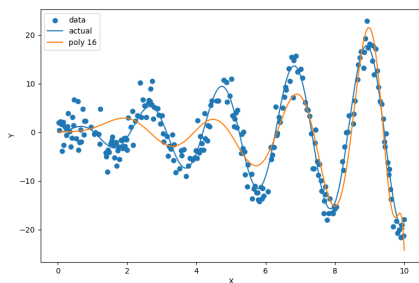
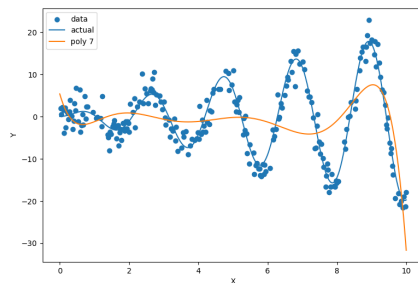
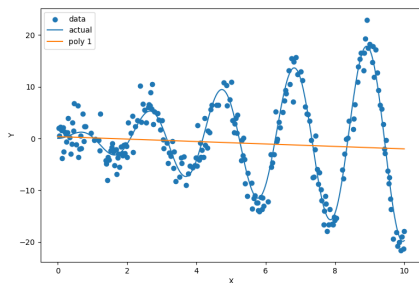
Non-linear transformations

► k-th order polynomial features on one variable

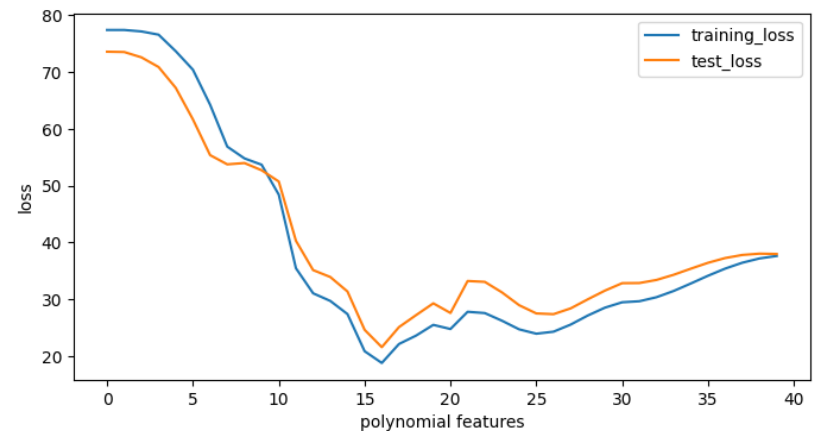
$$\mathbf{x} = [x] \quad \mathbf{z} = \begin{bmatrix} x^0 \\ x^1 \\ \vdots \\ x^k \end{bmatrix}$$

PolynomialFeatures(degree=4)

[[7]	[[1. 7. 49. 343. 2401.]
[[22]	[[1. 22. 484. 10648. 234256.]
[[-38]	[[1. -38. 1444. -54872. 2085136.]
[[-67]	[[1. -67. 4489. -300763. 20151121.]
[[0]]	[[1. 0. 0. 0. 0.]



Evaluating different degrees



Non-linear transformations

- ▶ k-th order polynomial features on two variables

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \mathbf{z} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1^2 \\ x_1x_2 \\ x_2^2 \end{bmatrix} \quad \text{degree } k=2$$

`PolynomialFeatures(degree=2)`

```
[[ 77  79]    [[ 1.  77.  79. 5929. 6083. 6241.]  
[-46 -23]    [ 1. -46. -23. 2116. 1058. 529.]  
[ 97 -87]    [ 1.  97. -87. 9409. -8439. 7569.]  
[ 53  26]    [ 1.  53.  26. 2809. 1378.  676.]  
[ 41 -8]]    [ 1.  41.  -8. 1681. -328.  64.]
```

Non-linear transformations

- ▶ k-th order polynomial features on d variables
 - “all polynomial combinations of the features with degree less than or equal to the specified degree”
- ▶ Issues:
 - be aware of computational cost
 - be aware of overfitting
- ▶ Alternatives?
 - use other non-linear functions: logarithmic, power, etc.

Data transformation

- ▶ Best Practices
 - transform after splitting data
 - fit transformers on **training data only**
 - handle missing values before transformation and use appropriate imputation
 - use pipelines to ensure reproducibility and to prevent data leakage

Embeddings

- ▶ Vector representations
 - embeddings transform discrete data (words, images, audio code) into continuous vector spaces, capturing semantic and syntactic meaning
 - key properties of embedding spaces:
 - similar items cluster together
 - meaningful vector arithmetic (e.g., king - man + woman ≈ queen)
 - fixed dimensionality (typically 256-4096 dimensions)
- ▶ Foundation models
 - large models trained on massive datasets
 - can produce high-quality embeddings that capture complex semantic and syntactic relationships
- ▶ Applications
 - similarity search, classification, generation, machine translation