

CSC 461: Machine Learning

Fall 2024


PyTorch / Autograd

Prof. Marco Alvarez, Computer Science
University of Rhode Island

Autograd

CPU vs GPU in practice

(CPU performance not well-optimized, a little unfair)



From CS231 @ Stanford

Computational Graphs


Numpy

```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)
```



From CS231 @ Stanford

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



From CS231 @ Stanford

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



Good:

Clean API, easy to write numeric code

Bad:

- Have to compute our own gradients
- Can't run on GPU

From CS231 @ Stanford

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch

N, D = 3, 4
x = torch.randn(N, D)
y = torch.randn(N, D)
z = torch.randn(N, D)

a = x * y
b = a + z
c = torch.sum(b)
```

Looks exactly like numpy!

From CS231 @ Stanford

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch

N, D = 3, 4
x = torch.randn(N, D, requires_grad=True)
y = torch.randn(N, D)
z = torch.randn(N, D)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

PyTorch handles gradients for us!

From CS231 @ Stanford

Automatic differentiation

Modern machine learning

- parameters (model weights) are adjusted according to the gradient of the loss function with respect to the given parameter

PyTorch has a built-in differentiation engine

- called **torch.autograd**
- supports automatic computation of gradient for any computational graph

Computational Graphs

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch
device = 'cuda:0'
N, D = 3, 4
x = torch.randn(N, D, requires_grad=True, device=device)
y = torch.randn(N, D, device=device)
z = torch.randn(N, D, device=device)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

Trivial to run on GPU - just construct arrays on a different device!


From CS231 @ Stanford

Example

Consider multinomial logistic regression

- input x , parameters w and b , and some loss function

```
x = torch.rand(5) # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = x @ w + b
loss = F.binary_cross_entropy_with_logits(z, y)
```



Example

Computing gradients

```
loss.backward()
print(w.grad)
print(b.grad)
```

```
tensor([[0.0360, 0.2630, 0.0477],
        [0.0218, 0.1590, 0.0289],
        [0.0367, 0.2680, 0.0486],
        [0.0181, 0.1320, 0.0240],
        [0.0221, 0.1613, 0.0293]])
tensor([0.0436, 0.3186, 0.0578])
```

Gradient tracking

▶ Tracking gradients

- by default all tensors with `requires_grad=True` are tracking their computational history and support gradient computation

▶ Stop tracking computations

- there are some cases when we do not need tracking
 - e.g., trained model and just want to perform inference, i.e. we only want to do forward computations
- advantages:
 - mark some parameters as **frozen parameters**
 - speed up computations when you are only doing forward pass (more efficient)

```
with torch.no_grad():  
    z = x @ w + b
```

Forward and backward computation

```
import torch  
  
x = torch.tensor([-1., -2.], requires_grad=True)  
w = torch.tensor([2., -3.], requires_grad=True)  
b = torch.tensor(-3., requires_grad=True)  
  
# forward pass  
f = 1 / (1 + torch.exp(-(w @ x + b)))  
  
# backward pass  
f.backward()  
  
print(w.grad, x.grad, b.grad)  
  
tensor([-0.1966, -0.3932]) tensor([ 0.3932, -0.5898]) tensor(0.1966)
```

Forward and backward computation

```
import torch  
  
x = torch.tensor([-1., -2.], requires_grad=True)  
w = torch.tensor([2., -3.], requires_grad=True)  
b = torch.tensor(-3., requires_grad=True)  
  
# forward pass  
f = torch.sigmoid(w @ x + b)  
  
# backward pass  
f.backward()  
  
print(w.grad, x.grad, b.grad)  
  
tensor([-0.1966, -0.3932]) tensor([ 0.3932, -0.5898]) tensor(0.1966)
```

Datasets and Data loaders

Working with datasets

PyTorch provides two data primitives

- `torch.utils.data.Dataset`
 - allows you to use pre-loaded datasets as well as your own data
 - stores the samples and their corresponding labels
- `torch.utils.data.DataLoader`
 - wraps an iterable around the Dataset to enable easy access to the samples

PyTorch domain libraries provide a number of pre-loaded datasets

- datasets subclass `torch.utils.data.Dataset` and implement functions specific to the particular data
- can be used to prototype and benchmark models
- include: Image Datasets, Text Datasets, and Audio Datasets

Example: FashionMNIST

- dataset of Zalando's article images consisting of 60,000 training examples and 10,000 test examples
- each example comprises a 28x28 grayscale image and an associated label from one of 10 classes

Forward and backward computation

```
from torchvision import datasets
from torchvision.transforms import ToTensor

training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor()
)

test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor()
)
```

Working with datasets

Datasets

- retrieves our dataset's features and labels one sample at a time

Dataloaders

- iterable that provides efficient mini-batch handling
- can reshuffle the data at every epoch and use Python's multiprocessing to speed up data retrieval

```
from torch.utils.data import DataLoader

train_dataloader = DataLoader(
    training_data,
    batch_size=64,
    shuffle=True)

X = torch.rand(20, 5)
loader = DataLoader(X, batch_size=5, shuffle=True)

test_dataloader = DataLoader(
    test_data,
    batch_size=64,
    shuffle=False)

for batch in loader:
    print(batch.shape)
```

Building models

Models comprise of layers/modules that perform operations on data

- `torch.nn` namespace provides all the needed building blocks
- every module subclasses the `nn.Module`
- every `nn.Module` subclass implements the operations on input data in the `forward` method

```
class MultinomialLR(nn.Module):
    def __init__(self, in_dim, out_dim):
        super().__init__()
        self.linear = nn.Linear(in_dim, out_dim)
        self.activation = nn.Softmax(dim=1)

    def forward(self, x):
        logits = self.linear(x)
        probs = self.activation(logits)
        return probs

model = MultinomialLR().to(device)
print(model)
```

Training/Testing a model

```
# set the model to training mode
model.train()
for batch in dataloader:
    pred = model(batch.x)
    loss = loss_fn(pred, batch.y)
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

# set the model to evaluation mode
model.eval()
test_loss = 0
with torch.no_grad():
    for batch in dataloader:
        pred = model(batch.x)
        test_loss += loss_fn(pred, y).item()
test_loss /= len(dataloader)
```